

# **Machine Learning Worksheet Solution 4**

Julius Jankowski

# 1 Least squares regression

## Problem 1:

See attached pages.

## Problem 2:

Starting from the error sum formulation:

$$E_{weighted} = \frac{1}{2} \sum_{i=1}^N t_i (\mathbf{w}^T \Phi(\mathbf{x}_i) - y_i)^2 \quad (1.1)$$

We can reformulate this problem by stacking all summation terms in a matrix vector style. The  $t_i$ s will be represented by an eye matrix  $\tilde{\mathbf{T}}$  with the corresponding  $t_i$  in row  $i$ :

$$E_{weighted} = \frac{1}{2} (\tilde{\Phi} \cdot \mathbf{w} - \mathbf{y})^T \cdot \tilde{\mathbf{T}} \cdot (\tilde{\Phi} \cdot \mathbf{w} - \mathbf{y}) \quad (1.2)$$

By resolving the brackets, we obtain ( $\tilde{\mathbf{T}}^T = \tilde{\mathbf{T}}$ ):

$$E_{weighted} = \frac{1}{2} (\mathbf{w}^T \cdot \tilde{\Phi}^T \cdot \tilde{\mathbf{T}} \cdot \tilde{\Phi} \cdot \mathbf{w} - 2\mathbf{y}^T \cdot \tilde{\mathbf{T}} \cdot \tilde{\Phi} \cdot \mathbf{w} + \mathbf{y}^T \cdot \tilde{\mathbf{T}} \cdot \mathbf{y}) \quad (1.3)$$

and now we can easily compute the derivative w.r.t. the weights  $\mathbf{w}$  in order to get the optimizer  $\mathbf{w}^*$ :

$$\frac{\partial E_{weighted}}{\partial \mathbf{w}} = \mathbf{w}^T \cdot \tilde{\Phi}^T \cdot \tilde{\mathbf{T}} \cdot \tilde{\Phi} - \mathbf{y}^T \cdot \tilde{\mathbf{T}} \cdot \tilde{\Phi} = 0 \quad (1.4)$$

$$\mathbf{w}^* = (\tilde{\Phi}^T \cdot \tilde{\mathbf{T}} \cdot \tilde{\Phi})^{-1} \cdot \tilde{\Phi}^T \cdot \tilde{\mathbf{T}} \cdot \mathbf{y} \quad (1.5)$$

Here we can see that there is a new version of a pseudo-inverse matrix  $\tilde{\Phi}^+ = (\tilde{\Phi}^T \cdot \tilde{\mathbf{T}} \cdot \tilde{\Phi})^{-1} \cdot \tilde{\Phi}^T \cdot \tilde{\mathbf{T}}$  which fits the general condition for an inverse matrix  $\tilde{\Phi}^+ \cdot \tilde{\Phi} = \tilde{\mathbf{I}}$ . Hence the  $t_i$ s can be interpreted as metric for the feature matrix  $\tilde{\Phi}$ , which weights each feature independently. From a machine learning perspective, the  $t_i$ s can be used for:

- 1) giving samples with less variance (hence more reliable) a higher impact in learning the  $\mathbf{w}^*$ .
- 2) replacing two exact copies of one sample with one sample with a weighting factor of 2.

## 2 Ridge regression

### Problem 3:

Since we can reformulate the error equation between sum style and matrix-vector style as follows:

$$E = \frac{1}{2}(\tilde{\Phi} \cdot \mathbf{w} - \mathbf{y})^T \cdot (\tilde{\Phi} \cdot \mathbf{w} - \mathbf{y}) = \frac{1}{2} \sum_{i=1}^N (\mathbf{w}^T \Phi_i - y_i)^2 \quad (2.1)$$

where  $\Phi_i$  corresponds to the  $i$ -th row vector of the design matrix. Hence we can also reformulate the specification stated in the task description:

$$E = \frac{1}{2} \sum_{i=1}^{N+M} (\mathbf{w}^T \Phi_i - y_i)^2 \quad (2.2)$$

and now split up the sum after  $N$  samples:

$$E = \frac{1}{2} \sum_{i=1}^N (\mathbf{w}^T \Phi_i - y_i)^2 + \frac{1}{2} \sum_{i=N+1}^{N+M} (\sqrt{\lambda} w_{i-N})^2 = \frac{1}{2} \sum_{i=1}^N (\mathbf{w}^T \Phi_i - y_i)^2 + \frac{\lambda}{2} \sum_{i=1}^M w_i^2 \quad (2.3)$$

which is equal to the ridge regression formulation.

## 3 Bayesian linear regression

Problem 4:

# 04\_homework\_linear\_regression

November 16, 2017

## 1 Programming assignment 4: Linear regression

```
In [5]: import numpy as np
```

```
from sklearn.datasets import load_boston
from sklearn.model_selection import train_test_split
```

### 1.1 Your task

In this notebook code skeleton for performing linear regression is given. Your task is to complete the functions where required. You are only allowed to use built-in Python functions, as well as any numpy functions. No other libraries / imports are allowed.

### 1.2 Load and preprocess the data

In this assignment we will work with the Boston Housing Dataset. The data consists of 506 samples. Each sample represents a district in the city of Boston and has 13 features, such as crime rate or taxation level. The regression target is the median house price in the given district (in \$1000's).

More details can be found here: <http://lib.stat.cmu.edu/datasets/boston>

```
In [6]: X , y = load_boston(return_X_y=True)
```

```
# Add a vector of ones to the data matrix to absorb the bias term
# (Recall slide #7 from the lecture)
X = np.hstack([np.ones([X.shape[0], 1]), X])
# From now on, D refers to the number of features in the AUGMENTED dataset (i.e. including the bias term)

# Split into train and test
test_size = 0.2
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size)
```

### 1.3 Task 1: Fit standard linear regression

```
In [28]: def fit_least_squares(X, y):
        """Fit ordinary least squares model to the data.

        Parameters
```

```

-----
X : array, shape [N, D]
    (Augmented) feature matrix.
y : array, shape [N]
    Regression targets.

Returns
-----
w : array, shape [D]
    Optimal regression coefficients (w[0] is the bias term).

"""

w = np.linalg.inv(X.transpose().dot(X)).dot(X.transpose()).dot(y)

return w

```

## 1.4 Task 2: Fit ridge regression

```

In [32]: def fit_ridge(X, y, reg_strength):
    """Fit ridge regression model to the data.

    Parameters
    -----
    X : array, shape [N, D]
        (Augmented) feature matrix.
    y : array, shape [N]
        Regression targets.
    reg_strength : float
        L2 regularization strength (denoted by lambda in the lecture)

    Returns
    -----
    w : array, shape [D]
        Optimal regression coefficients (w[0] is the bias term).

    """

    D = X.shape[1]
    lambda_matrix = reg_strength*np.eye(D)
    w = np.linalg.inv(lambda_matrix + X.transpose().dot(X)).dot(X.transpose()).dot(y)

    return w

```

## 1.5 Task 3: Generate predictions for new data

```

In [25]: def predict_linear_model(X, w):
    """Generate predictions for the given samples.

```

```

Parameters
-----
X : array, shape [N, D]
    (Augmented) feature matrix.
w : array, shape [D]
    Regression coefficients.

Returns
-----
y_pred : array, shape [N]
    Predicted regression targets for the input data.
"""

y_pred = X.dot(w)

return y_pred

```

## 1.6 Task 4: Mean squared error

```

In [30]: def mean_squared_error(y_true, y_pred):
    """Compute mean squared error between true and predicted regression targets.

    Reference: `https://en.wikipedia.org/wiki/Mean_squared_error`

    Parameters
    -----
    y_true : array
        True regression targets.
    y_pred : array
        Predicted regression targets.

    Returns
    -----
    mse : float
        Mean squared error.

    """

    n = len(y_true)
    E = 0
    for i in range(n):
        E = E + np.dot((y_true[i] - y_pred[i]), (y_true[i] - y_pred[i]))

    return E / n

```

## 1.7 Compare the two models

The reference implementation produces \* MSE for Least squares  $\approx 23.98$  \* MSE for Ridge regression  $\approx 21.05$

Your results might be slightly (i.e.  $\pm 1\%$ ) different from the reference solution due to numerical reasons.

```
In [33]: # Load the data
np.random.seed(1234)
X, y = load_boston(return_X_y=True)
X = np.hstack([np.ones([X.shape[0], 1]), X])
test_size = 0.2
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=test_size)

# Ordinary least squares regression
w_ls = fit_least_squares(X_train, y_train)
y_pred_ls = predict_linear_model(X_test, w_ls)
mse_ls = mean_squared_error(y_test, y_pred_ls)
print('MSE for Least squares = {}'.format(mse_ls))

# Ridge regression
reg_strength = 1
w_ridge = fit_ridge(X_train, y_train, reg_strength)
y_pred_ridge = predict_linear_model(X_test, w_ridge)
mse_ridge = mean_squared_error(y_test, y_pred_ridge)
print('MSE for Ridge regression = {}'.format(mse_ridge))

MSE for Least squares = 23.984307611774053
MSE for Ridge regression = 21.051487033771537
```